

Model Interpretation

Objectives

- Partial Dependence Plots
- Shapley Values

Pre-reads

1. Kaggle / Dan Becker: Machine Learning Explainability
 - <https://www.kaggle.com/dansbecker/partial-plots> (<https://www.kaggle.com/dansbecker/partial-plots>)
 - <https://www.kaggle.com/dansbecker/shap-values> (<https://www.kaggle.com/dansbecker/shap-values>)
2. Christoph Molnar: Interpretable Machine Learning
 - <https://christophm.github.io/interpretable-ml-book/pdp.html> (<https://christophm.github.io/interpretable-ml-book/pdp.html>)
 - <https://christophm.github.io/interpretable-ml-book/shapley.html> (<https://christophm.github.io/interpretable-ml-book/shapley.html>)

Libraries

- [PDPbox](https://github.com/SauceCat/PDPbox) (<https://github.com/SauceCat/PDPbox>): `pip install pdpbox`
- [shap](https://github.com/slundberg/shap) (<https://github.com/slundberg/shap>): `conda install -c conda-forge shap / pip install shap`

Types of explanations

Global explanation: all features in relation to each other

- Feature Importances (mean decrease impurity)
- Permutation Importances

- Drop-Column Importances

Global explanation: individual feature in relation to target

- Partial Dependence plots

Individual prediction explanation

- Shapley Values

Note that the coefficients from a linear model give you all three types of explanations!

Titanic

```
In [4]: !echo y | conda install -c conda-forge seaborn
```

Collecting package metadata (repodata.json): done
Solving environment: done

Package Plan

environment location: /Library/anaconda3

added / updated specs:
- seaborn

The following packages will be downloaded:

package	build		
seaborn-0.9.0	py_1	163 KB	conda-forge
Total:		163 KB	

The following NEW packages will be INSTALLED:

seaborn conda-forge/noarch::seaborn-0.9.0-py_1

Proceed ([y]/n)?

Downloading and Extracting Packages
seaborn-0.9.0 | 163 KB | ##### | 100%
Preparing transaction: done
Verifying transaction: done
Executing transaction: done

```
In [5]: %matplotlib inline
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_score

def load_titanic():
    df = sns.load_dataset('titanic')
    df['age'] = df['age'].fillna(df['age'].mean())
    df['class'] = df['class'].map({'First': 1, 'Second': 2, 'Third': 3})
    df['female'] = df['sex'] == 'female'
    X = df[['age', 'class', 'fare', 'female']]
    y = df['survived']
    return X, y

X, y = load_titanic()
```

Naive majority class baseline

```
In [6]: y.value_counts(normalize=True)
```

```
Out[6]: 0    0.616162
        1    0.383838
        Name: survived, dtype: float64
```

Logistic Regression

```
In [7]: lr = LogisticRegression(solver='lbfgs')  
cross_val_score(lr, X, y, scoring='accuracy', cv=5, n_jobs=-1)
```

```
Out[7]: array([0.7877095 , 0.78212291, 0.78651685, 0.7752809 , 0.80225989])
```

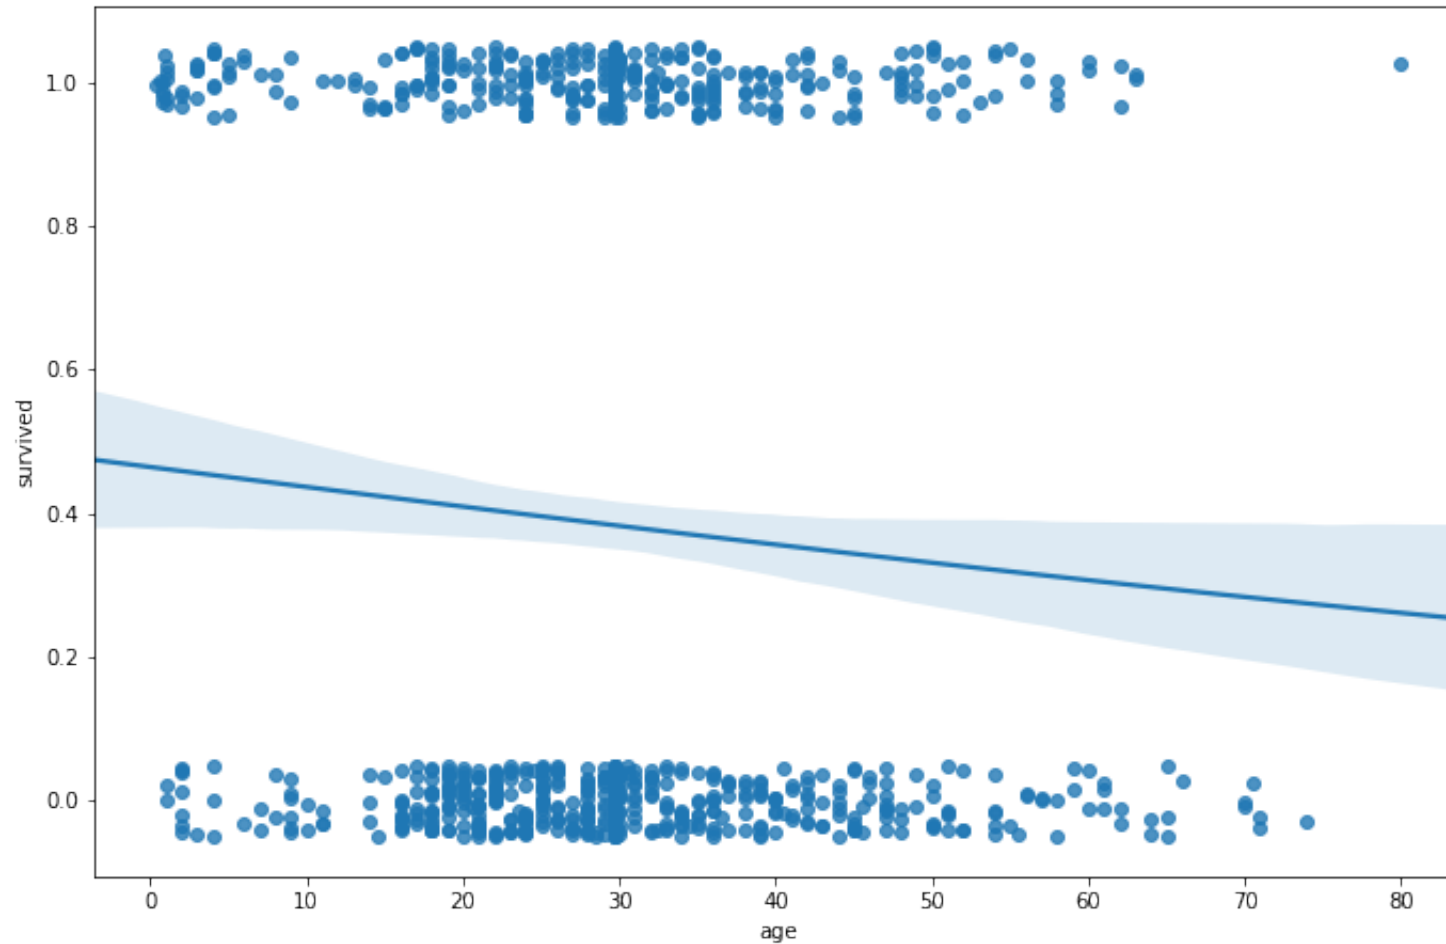
```
In [8]: lr.fit(X, y)  
pd.Series(lr.coef_[0], X.columns)
```

```
Out[8]: age      -0.032595  
class    -1.112937  
fare       0.000805  
female     2.512794  
dtype: float64
```

```
In [9]: g = sns.regplot(x=X['age'], y=y, logistic=True, y_jitter=.05)  
g.figure.set_size_inches(11.5, 7.5)
```

/Library/anaconda3/lib/python3.7/site-packages/scipy/stats/stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```



Gradient Boosting

```
In [10]: gb = GradientBoostingClassifier()
cross_val_score(gb, X, y, scoring='accuracy', cv=5, n_jobs=-1)
```

Out[10]: array([0.80446927, 0.82681564, 0.85393258, 0.83146067, 0.84745763])

```
In [11]: gb.fit(X, y)
pd.Series(gb.feature_importances_, X.columns)
```

Out[11]: age 0.131213
class 0.173489
fare 0.199991
female 0.495308
dtype: float64

```
In [20]: !echo y | conda update -n base -c defaults conda
```

Collecting package metadata (repodata.json): done
Solving environment: done

Package Plan

environment location: /Library/anaconda3

added / updated specs:
- conda

The following packages will be downloaded:

package	build	
-----	-----	
_py-xgboost-mutex-2.0	cpu_0	8 KB

astor-0.8.0	py37_0	45 KB
c-ares-1.15.0	h1de35cc_1001	83 KB
clang_osx-64-4.0.1	h1ce6c1d_16	140 KB
clangxx_osx-64-4.0.1	h22b1bf0_16	140 KB
conda-4.7.10	py37_0	3.0 MB
conda-package-handling-1.3.11	py37_0	260 KB
dill-0.3.0	py37_0	116 KB
glue-core-0.14.2	py37_0	1.4 MB
libprotobuf-3.8.0	hd9629dc_0	4.4 MB
markdown-3.1.1	py37_0	113 KB
mock-3.0.5	py37_0	47 KB
mpl-scatter-density-0.6	py_0	647 KB
plotly-4.0.0	py_0	3.8 MB
protobuf-3.8.0	py37h0a44026_0	678 KB
python-graphviz-0.10.1	py_0	22 KB
tabulate-0.8.3	py37_0	38 KB

Total:		14.9 MB

The following packages will be REMOVED:

pbr-5.1.3-py_0

The following packages will be UPDATED:

astor	0.7.1-py37_0 --> 0.8.0-py37_0
c-ares	1.15.0-h1de35cc_1 --> 1.15.0-h1de35cc_1001
clang_osx-64	4.0.1-h1ce6c1d_11 --> 4.0.1-h1ce6c1d_16
clangxx_osx-64	4.0.1-h22b1bf0_11 --> 4.0.1-h22b1bf0_16
dill	0.2.9-py37_0 --> 0.3.0-py37_0
glue-core	0.14.1-py37_0 --> 0.14.2-py37_0
libprotobuf	3.7.1-hd9629dc_0 --> 3.8.0-hd9629dc_0
markdown	3.1-py37_0 --> 3.1.1-py37_0
mock	2.0.0-py37_0 --> 3.0.5-py37_0
mpl-scatter-densi~	0.5-py_0 --> 0.6-py_0
plotly	3.7.0-py_0 --> 4.0.0-py_0
protobuf	3.7.1-py37h0a44026_0 --> 3.8.0-py37h0a44026_0

The following packages will be SUPERSEDED by a higher-priority channel:

```
_py-xgboost-mutex          conda-forge --> pkgs/main
conda                      conda-forge --> pkgs/main
conda-package-han~ conda-forge::conda-package-handling-1~ --> pkgs/main::conda-package-handling-1.3.11-py37_0
python-graphviz            conda-forge --> pkgs/main
tabulate                   conda-forge/noarch::tabulate-0.8.3-py~ --> pkgs/main/osx-64::tabulate-0.8.3-py37_0
```

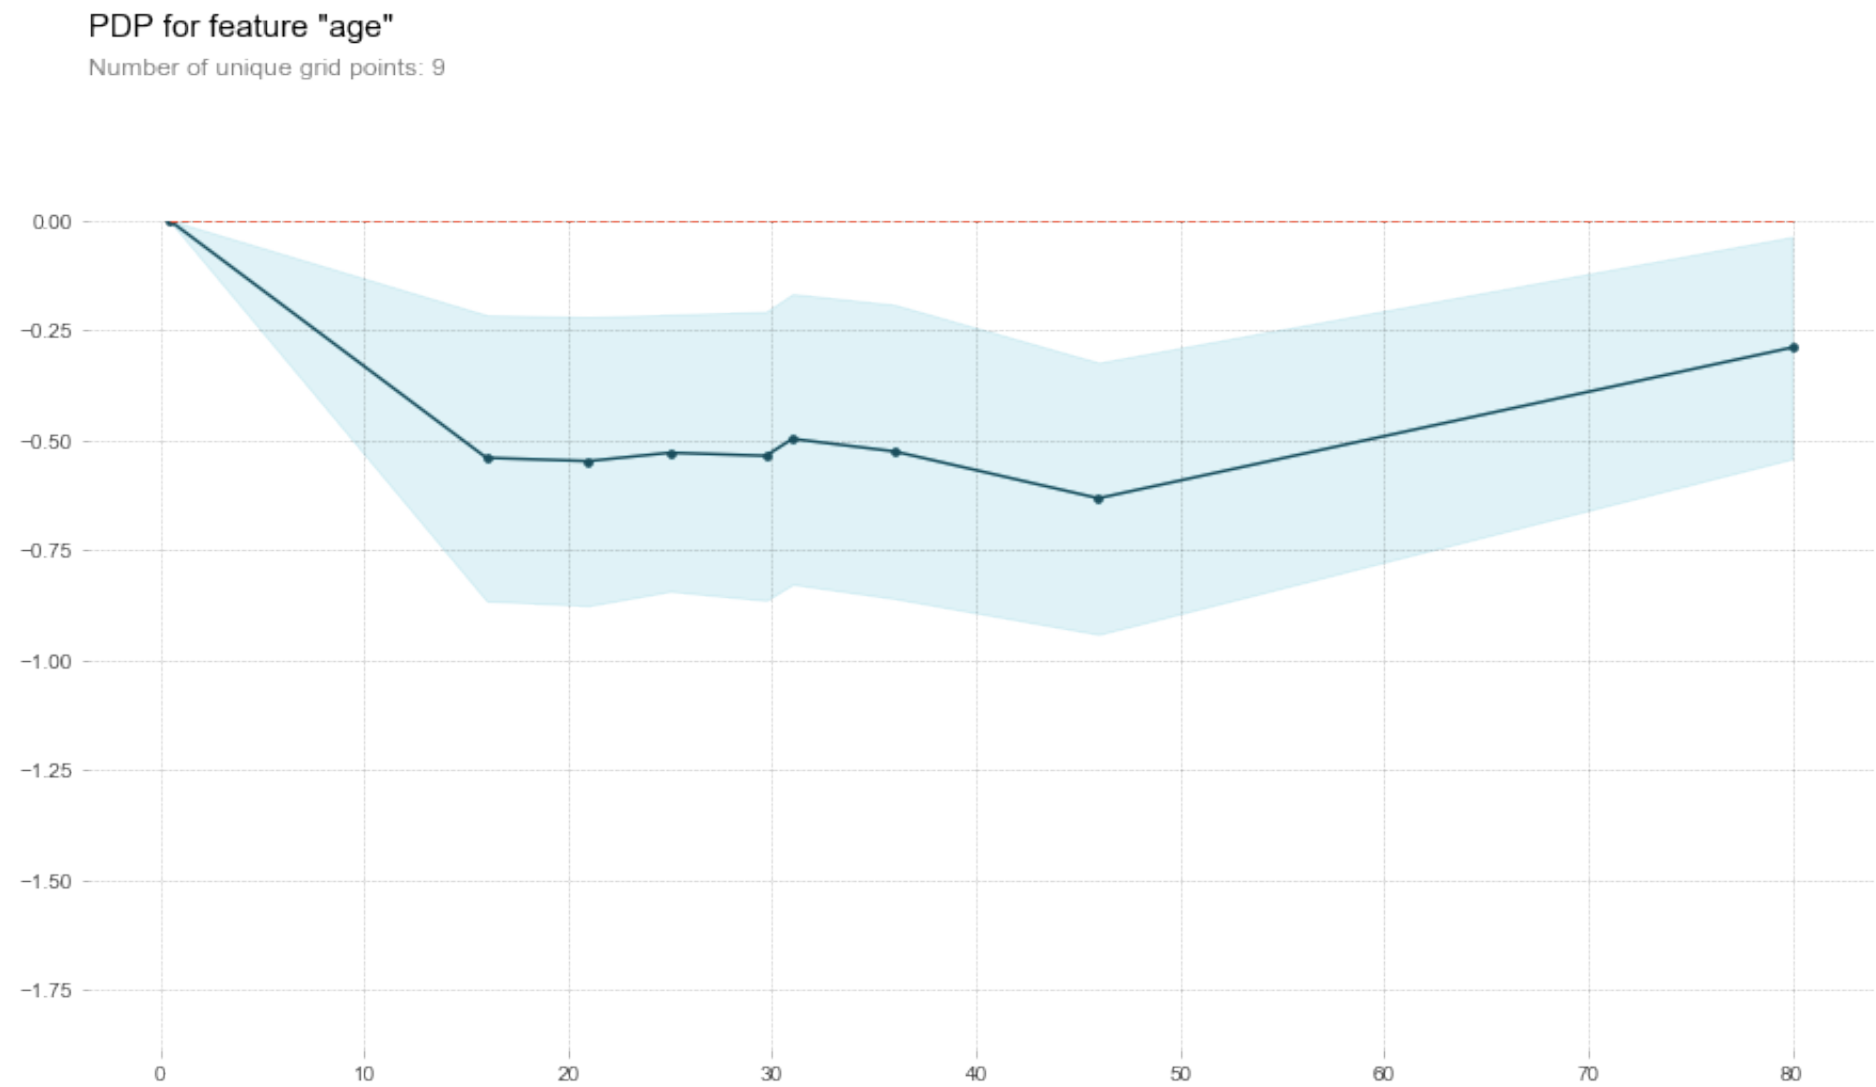
Proceed ([y]/n)?

```
Downloading and Extracting Packages
dill-0.3.0 | 116 KB | ##### | 100%
protobuf-3.8.0 | 678 KB | ##### | 100%
python-graphviz-0.10 | 22 KB | ##### | 100%
conda-4.7.10 | 3.0 MB | ##### | 100%
astor-0.8.0 | 45 KB | ##### | 100%
markdown-3.1.1 | 113 KB | ##### | 100%
libprotobuf-3.8.0 | 4.4 MB | ##### | 100%
tabulate-0.8.3 | 38 KB | ##### | 100%
clang_osx-64-4.0.1 | 140 KB | ##### | 100%
conda-package-handli | 260 KB | ##### | 100%
plotly-4.0.0 | 3.8 MB | ##### | 100%
clangxx_osx-64-4.0.1 | 140 KB | ##### | 100%
mpl-scatter-density- | 647 KB | ##### | 100%
c-ares-1.15.0 | 83 KB | ##### | 100%
glue-core-0.14.2 | 1.4 MB | ##### | 100%
mock-3.0.5 | 47 KB | ##### | 100%
_py-xgboost-mutex-2. | 8 KB | ##### | 100%
Preparing transaction: done
Verifying transaction: done
Executing transaction: done
```

```
In [21]: from pdpbox.pdp import pdp_isolate, pdp_plot
feature='age'
```

```
feature = age
pdp_isolated = pdp_isolate(model=gb, dataset=X, model_features=X.columns,
                           feature=feature)

pdp_plot(pdp_isolated, feature);
```



age

From [PDPbox documentation \(https://pdpbox.readthedocs.io/en/latest/\)](https://pdpbox.readthedocs.io/en/latest/):

The common headache: When using black box machine learning algorithms like random forest and boosting, it is hard to understand the relations between predictors and model outcome. For example, in terms of random forest, all we get is the feature importance. Although we can know which feature is significantly influencing the outcome based on the importance calculation, it really sucks that we don't know in which direction it is influencing. And in most of the real cases, the effect is non-monotonic. We need some powerful tools to help understanding the complex relations between predictors and model prediction.

```
In [18]: !echo y | conda install -c conda-forge pdpbox
```

```
Collecting package metadata: done  
Solving environment: done
```

```
==> WARNING: A newer version of conda exists. <==  
current version: 4.6.14  
latest version: 4.7.10
```

```
Please update conda by running
```

```
$ conda update -n base -c defaults conda
```

```
## Package Plan ##
```

```
environment location: /Library/anaconda3
```

```
added / updated specs:  
- pdpbox
```

The following packages will be downloaded:

package	build		
-----	-----		
_r-mutex-1.0.1	anacondar_1	3 KB	conda-forge
conda-4.7.10	py37_0	3.0 MB	conda-forge
conda-package-handling-1.4.1	py37_0	261 KB	conda-forge
matplotlib-base-3.1.1	py37h3a684a6_1	6.6 MB	conda-forge
pdpbox-0.2.0	py_0	55.1 MB	conda-forge
-----	-----		
Total:		65.0 MB	

The following NEW packages will be INSTALLED:

```
conda-package-han~ conda-forge/osx-64::conda-package-handling-1.4.1-py37_0
matplotlib-base    conda-forge/osx-64::matplotlib-base-3.1.1-py37h3a684a6_1
pdpbox             conda-forge/noarch::pdpbox-0.2.0-py_0
```

The following packages will be UPDATED:

```
_r-mutex          pkgs/r/osx-64::_r-mutex-1.0.0-anacond~ --> conda-forge/noarch::_r-mutex-1.0.1-anacondar_1
conda              pkgs/main::conda-4.6.14-py37_0 --> conda-forge::conda-4.7.10-py37_0
```

Proceed ([y]/n)?

Downloading and Extracting Packages

conda-package-handli	261 KB	#####	100%
matplotlib-base-3.1.	6.6 MB	#####	100%
conda-4.7.10	3.0 MB	#####	100%
pdpbox-0.2.0	55.1 MB	#####	100%
_r-mutex-1.0.1	3 KB	#####	100%

Preparing transaction: done
Verifying transaction: done
Executing transaction: done

Animation by Christoph Molnar (<https://twitter.com/ChristophMolnar/status/1066398522608635904>), author of *Interpretable Machine Learning* (<https://christophm.github.io/interpretable-ml-book/>).

Partial dependence plots show how a feature affects predictions of a Machine Learning model on average.

1. Define grid along feature
2. Model predictions at grid points
3. Line per data instance -> ICE (Individual Conditional Expectation) curve
4. Average curves to get a PDP (Partial Dependence Plot)

Compare Predictions

```
In [16]: from sklearn.model_selection import cross_val_predict

y_pred_lr = cross_val_predict(lr, X, y, cv=5, n_jobs=-1)
y_pred_gb = cross_val_predict.gb, X, y, cv=5, n_jobs=-1)

preds = pd.DataFrame({'true': y, 'lr': y_pred_lr, 'gb': y_pred_gb})

gb_right = preds['gb'] == preds['true']
lr_wrong = preds['lr'] != preds['true']

len(preds[gb_right & lr_wrong]) / len(preds)
```

Out[16]: 0.08417508417508418

```
In [17]: preds[gb_right & lr_wrong].head()
```

Out[17]:

	true	lr	gb
27	0	1	0
64	0	1	0
71	0	1	0
74	1	0	1
78	1	0	1

```
In [18]: data_for_prediction = X.loc[27]

data_for_prediction
```

Out[18]: age 19
class 1
fare 263
female False
Name: 27, dtype: object

Explain individual prediction

<https://www.kaggle.com/dansbecker/shap-values> (<https://www.kaggle.com/dansbecker/shap-values>)

```
In [30]: #pip install shap  
!echo y | conda install -c conda-forge shap
```

```
Collecting package metadata (repodata.json): done  
Solving environment: done
```

```
## Package Plan ##
```

```
environment location: /Library/anaconda3
```

```
added / updated specs:  
- shap
```

```
The following NEW packages will be INSTALLED:
```

```
shap                      conda-forge/osx-64::shap-0.29.3-py37h86efe34_0
```


```
Proceed ([y]/n)?  
Preparing transaction: done  
Verifying transaction: done  
Executing transaction: done
```

```
In [19]: import shap

# Create object that can calculate shap values
explainer = shap.TreeExplainer(gb)

# Calculate Shap values
shap_values = explainer.shap_values(data_for_prediction)

shap.initjs()
shap.force_plot(explainer.expected_value, shap_values, data_for_prediction)
```



Lending Club

```
In [20]: import category_encoders as ce
import pandas as pd
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.model_selection import cross_val_score
from sklearn.pipeline import make_pipeline

# Load data from https://www.kaggle.com/c/dsl-tree-ensembles/data
X_train = pd.read_csv('/Users/wel51x/Downloads/ds2-tree-ensembles/train_features.csv')
X_test = pd.read_csv('/Users/wel51x/Downloads/ds2-tree-ensembles/test_features.csv')
y_train = pd.read_csv('/Users/wel51x/Downloads/ds2-tree-ensembles/train_labels.csv')['charged_off']
```



```

train = pd.read_csv('/Users/wel51x/Downloads/ds2-tree-ensembles/train_labels.csv', ['Charged_Off'])
sample_submission = pd.read_csv('/Users/wel51x/Downloads/ds2-tree-ensembles/sample_submission.csv')

def wrangle(X):
    X = X.copy()

    # Drop some columns
    X = X.drop(columns='id') # id is random
    X = X.drop(columns=['member_id', 'url', 'desc']) # All null
    X = X.drop(columns='title') # Duplicative of purpose
    X = X.drop(columns='grade') # Duplicative of sub_grade

    # Transform sub_grade from "A1" - "G5" to 1.1 - 7.5
    def wrangle_sub_grade(x):
        first_digit = ord(x[0]) - 64
        second_digit = int(x[1])
        return first_digit + second_digit/10

    X['sub_grade'] = X['sub_grade'].apply(wrangle_sub_grade)

    # Convert percentages from strings to floats
    X['int_rate'] = X['int_rate'].str.strip('%').astype(float)
    X['revol_util'] = X['revol_util'].str.strip('%').astype(float)

    # Transform earliest_cr_line to an integer: how many days it's been open
    X['earliest_cr_line'] = pd.to_datetime(X['earliest_cr_line'], infer_datetime_format=True)
    X['earliest_cr_line'] = pd.Timestamp.today() - X['earliest_cr_line']
    X['earliest_cr_line'] = X['earliest_cr_line'].dt.days

    # Create features for three employee titles: teacher, manager, owner
    X['emp_title'] = X['emp_title'].str.lower()
    X['emp_title_teacher'] = X['emp_title'].str.contains('teacher', na=False)
    X['emp_title_manager'] = X['emp_title'].str.contains('manager', na=False)
    X['emp_title_owner'] = X['emp_title'].str.contains('owner', na=False)

    # Drop categoricals with high cardinality
    X = X.drop(columns=['emp_title', 'zip code'])

```

```
# Transform features with many nulls to binary flags
many_nulls = ['sec_app_mths_since_last_major_derog',
              'sec_app_revol_util',
              'sec_app_earliest_cr_line',
              'sec_app_mort_acc',
              'dti_joint',
              'sec_app_collections_12_mths_ex_med',
              'sec_app_chargeoff_within_12_mths',
              'sec_app_num_rev_accts',
              'sec_app_open_act_il',
              'sec_app_open_acc',
              'revol_bal_joint',
              'annual_inc_joint',
              'sec_app_inq_last_6mths',
              'mths_since_last_record',
              'mths_since_recent_bc_dlq',
              'mths_since_last_major_derog',
              'mths_since_recent_revol_delinq',
              'mths_since_last_delinq',
              'il_util',
              'emp_length',
              'mths_since_recent_inq',
              'mo_sin_old_il_acct',
              'mths_since_rcnt_il',
              'num_tl_120dpd_2m',
              'bc_util',
              'percent_bc_gt_75',
              'bc_open_to_buy',
              'mths_since_recent_bc']

for col in many_nulls:
    X[col] = X[col].isnull()

# For features with few nulls, do mean imputation
for col in X:
    if X[col].isnull().sum() > 0:
```

```

        X[col] = X[col].fillna(X[col].mean())

    # Return the wrangled dataframe
    return X

# Wrangle train and test in the same way
X_train = wrangle(X_train)
X_test  = wrangle(X_test)

```

```

In [25]: %%time
from sklearn.metrics import roc_auc_score
from sklearn.model_selection import train_test_split

X_train, X_val, y_train, y_val = train_test_split(
    X_train, y_train, test_size=0.2, stratify=y_train, random_state=42)

encoder = ce.OrdinalEncoder()
X_train = encoder.fit_transform(X_train)
X_val = encoder.transform(X_val)
gb = GradientBoostingClassifier()
gb.fit(X_train, y_train)
y_pred_proba = gb.predict_proba(X_val)[:,-1]
print('Validation ROC AUC:', roc_auc_score(y_val, y_pred_proba))

```

Validation ROC AUC: 0.7226505704278647
 CPU times: user 14min 17s, sys: 14.4 s, total: 14min 31s
 Wall time: 15min 2s

Partial Dependence Plot

```


In [26]: from pdpbox.pdp import pdp_isolate, pdp_plot

feature='int_rate'
pdp_isolated = pdp_isolate(model=gb, dataset=X_val,

```

```
model_features=X_val.columns, feature=feature)
pdp_plot(pdp_isolated, feature);
```

PDP for feature "int_rate"
Number of unique grid points: 10



Individual predictions

```
In [27]: import numpy as np
y_pred = (y_pred_proba >= 0.5).astype(int)
confidence = np.abs(y_pred_proba - 0.5)
preds = pd.DataFrame({'y_val': y_val, 'y_pred': y_pred,
                      'y_pred_proba': y_pred_proba,
                      'confidence': confidence})

preds.head()
```

Out[27]:

	y_val	y_pred	y_pred_proba	confidence
734948	0	0	0.120889	0.379111
662775	1	0	0.444685	0.055315
1293995	0	0	0.184529	0.315471
121846	0	0	0.055538	0.444462
1040519	0	0	0.105202	0.394798

```
In [28]: # True positives, with high confidence
preds[(y_val==1) & (y_pred==1)].sort_values(by='confidence', ascending=False).head()
```


Out[28]:

	y_val	y_pred	y_pred_proba	confidence
509674	1	1	0.778066	0.278066
1075507	1	1	0.746171	0.246171
1237423	1	1	0.729384	0.229384
525726	1	1	0.725751	0.225751
411364	1	1	0.725219	0.225219

```
In [30]: X_val.shape
```

Out[30]: (209513, 98)

```
In [32]: data_for_prediction = X_val.iloc[17575]
explainer = shap.TreeExplainer(gb)
shap_values = explainer.shap_values(data_for_prediction)
shap.force_plot(explainer.expected_value, shap_values, data_for_prediction)
```




```
In [33]: # True negatives, with high confidence
preds[(y_val==0) & (y_pred==0)].sort_values(by='confidence', ascending=False).head()
```

Out[33]:

	y_val	y_pred	y_pred_proba	confidence
716929	0	0	0.025035	0.474965
1024450	0	0	0.025919	0.474081
440926	0	0	0.025988	0.474012
626127	0	0	0.026001	0.473999
1129419	0	0	0.026023	0.473977

```
In [34]: data_for_prediction = X_val.loc[1778]
shap_values = explainer.shap_values(data_for_prediction)
shap.force_plot(explainer.expected_value, shap_values, data_for_prediction)
```




```
In [35]: # False positives, with high (mistaken) confidence
preds[(y_val==0) & (y_pred==1)].sort_values(by='confidence', ascending=False).head()
```

Out[35]:

	y_val	y_pred	y_pred_proba	confidence
929700	0	1	0.733019	0.233019
1124876	0	1	0.724849	0.224849
1218998	0	1	0.713815	0.213815
14044	0	1	0.708883	0.208883
314029	0	1	0.708107	0.208107

```
In [36]: data_for_prediction = X_val.loc[33542]
shap_values = explainer.shap_values(data_for_prediction)
shap.force_plot(explainer.expected_value, shap_values, data_for_prediction)
```



```
In [37]: data_for_prediction
```

Out[37]:

loan_amnt	12000
funded_amnt	12000
term	1

int_rate	8.18
installment	377.04
sub_grade	2.1
emp_length	False
home_ownership	3
annual_inc	100000
purpose	2
addr_state	9
dti	5.09
delinq_2yrs	0
earliest_cr_line	7035
inq_last_6mths	0
mths_since_last_delinq	False
mths_since_last_record	False
open_acc	7
pub_rec	1
revol_bal	15692
revol_util	61.3
total_acc	18
initial_list_status	2
collections_12_mths_ex_med	0
mths_since_last_major_derog	True
application_type	1
annual_inc_joint	True
dti_joint	True
acc_now_delinq	0
tot_coll_amt	0
...	
num_rev_accts	11
num_rev_tl_bal_gt_0	5
num_sats	7
num_tl_120dpd_2m	False
num_tl_30dpd	0
num_tl_90g_dpd_24m	0
num_tl_op_past_12m	0
pct_tl_nvr_dlq	94.4
percent_bc_gt_75	False


```
pub_rec_bankruptcies      0
tax_liens                  1
tot_hi_cred_lim           25600
total_bal_ex_mort         15692
total_bc_limit            21900
total_il_high_credit_limit 0
revol_bal_joint           True
sec_app_earliest_cr_line   True
sec_app_inq_last_6mths     True
sec_app_mort_acc           True
sec_app_open_acc           True
sec_app_revol_util         True
sec_app_open_act_il        True
sec_app_num_rev_accts      True
sec_app_chargeoff_within_12_mths True
sec_app_collections_12_mths_ex_med True
sec_app_mths_since_last_major_derog True
disbursement_method       1
emp_title_teacher          False
emp_title_manager          True
emp_title_owner            False
Name: 33542, Length: 98, dtype: object
```

```
In [38]: # False negatives, with high (mistaken) confidence
preds[(y_val==1) & (y_pred==0)].sort_values(by='confidence', ascending=False).head()
```

Out[38]:

	y_val	y_pred	y_pred_proba	confidence
923003	1	0	0.028394	0.471606
374806	1	0	0.029279	0.470721
415599	1	0	0.031176	0.468824
1117287	1	0	0.031621	0.468379
1049990	1	0	0.031804	0.468196

```
In [40]: data_for_prediction = X_val.iloc[30492]
shap_values = explainer.shap_values(data_for_prediction)
shap.force_plot(explainer.expected_value, shap_values, data_for_prediction)
```




```
In [41]: # Most uncertain predictions (least confidence)
preds.sort_values(by='confidence', ascending=True).head()
```

Out[41]:

	y_val	y_pred	y_pred_proba	confidence
1224267	0	0	0.499996	0.000004
1041474	0	0	0.499991	0.000009
1115932	0	0	0.499979	0.000021
926988	1	0	0.499973	0.000027
1032678	0	0	0.499969	0.000031

```
In [43]: data_for_prediction = X_val.iloc[33095]
shap_values = explainer.shap_values(data_for_prediction)
shap.force_plot(explainer.expected_value, shap_values,
                data_for_prediction)
```



```
In [44]: data_for_prediction
```

Out[44]:

loan_amnt	15600
funded_amnt	15600
term	1
int_rate	26.3

installment	631.03
sub_grade	5.5
emp_length	False
home_ownership	1
annual_inc	42000
purpose	2
addr_state	32
dti	30.66
delinq_2yrs	1
earliest_cr_line	6152
inq_last_6mths	0
mths_since_last_delinq	False
mths_since_last_record	True
open_acc	11
pub_rec	0
revol_bal	5187
revol_util	66.5
total_acc	21
initial_list_status	2
collections_12_mths_ex_med	0
mths_since_last_major_derog	False
application_type	1
annual_inc_joint	True
dti_joint	True
acc_now_delinq	0
tot_coll_amt	0
...	
num_rev_accts	10
num_rev_tl_bal_gt_0	5
num_sats	11
num_tl_120dpd_2m	False
num_tl_30dpd	0
num_tl_90g_dpd_24m	0
num_tl_op_past_12m	3
pct_tl_nvr_dlq	76.2
percent_bc_gt_75	False
pub_rec_bankruptcies	0

```
tax_liens                                0
tot_hi_cred_lim                         164158
total_bal_ex_mort                       39247
total_bc_limit                          2500
total_il_high_credit_limit              43257
revol_bal_joint                         True
sec_app_earliest_cr_line                True
sec_app_inq_last_6mths                  True
sec_app_mort_acc                        True
sec_app_open_acc                        True
sec_app_revol_util                       True
sec_app_open_act_il                     True
sec_app_num_rev_accts                   True
sec_app_chargeoff_within_12_mths        True
sec_app_collections_12_mths_ex_med      True
sec_app_mths_since_last_major_derog     True
disbursement_method                     1
emp_title_teacher                       False
emp_title_manager                       False
emp_title_owner                         False
Name: 329957, Length: 98, dtype: object
```

In [0]: